

Identifying Minimal Changes in the Zone Abstract Domain

Kenny Ballou
Elena Sherman

Boise State University
Boise, Idaho
United States of America

July 2023



- Thank you.
- I'm Kenny Ballou and I will be presenting our work on Identifying Minimal Changes in the Zone Abstract Domain.

Outline

① Background and Motivation

- Zones Domain

- Exploiting DFA Features

② Algorithms and Approach

- Spurious Connections

- Connected Components

- Node Neighbors

- Minimal Neighbors

③ Experimental Results

- Application

④ Conclusions

Static analysis computes invariants

- Static analysis computes facts about programs. These facts can be represented in various ways.
- A common approach for numerical abstract interpretation, is a restricted set of inequalities

Static analysis computes invariants

Unit difference, two-variables per inequality

$$x - Z_0 = 0$$

$$w - x \leq 2$$

- Static analysis computes facts about programs. These facts can be represented in various ways.
- A common approach for numerical abstract interpretation, is a restricted set of inequalities

Static analysis computes invariants

Inequalities as invariants for a simple program

```
1 int example(int w, int y) {  
2     int x = 0;  
3     if (w <= x + 2) {  
4         if (y <= x) {  
5             assert y <= 0;  
6         }  
7     }  
8     return x;  
9 }
```

- Given some code, we produce a CFG
- Using DFA, we compute some invariants, for example
 - These two trapezoids represent the in-state and the out-state for the inner if
- Let's focus on this invariant, and describe briefly how Zones work


Static analysis computes invariants

Inequalities as invariants for a simple program

```

1 int example(int w, int y) {
2   int x = 0;
3   if (w <= x + 2) {
4     if (y <= x) {
5       assert y <= 0;
6     }
7   }
8   return x;
9 }

```



- Given some code, we produce a CFG
- Using DFA, we compute some invariants, for example
 - These two trapezoids represent the in-state and the out-state for the inner if
- Let's focus on this invariant, and describe briefly how Zones work

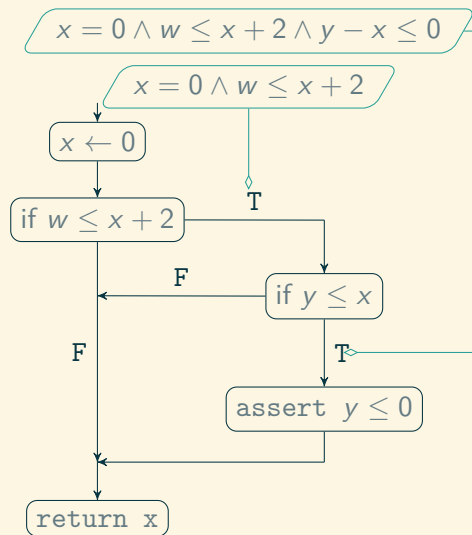
Static analysis computes invariants

Inequalities as invariants for a simple program

```

1 int example(int w, int y) {
2   int x = 0;
3   if (w <= x + 2) {
4     if (y <= x) {
5       assert y <= 0;
6     }
7   }
8   return x;
9 }

```



- Given some code, we produce a CFG
- Using DFA, we compute some invariants, for example
 - These two trapezoids represent the in-state and the out-state for the inner if
- Let's focus on this invariant, and describe briefly how Zones work

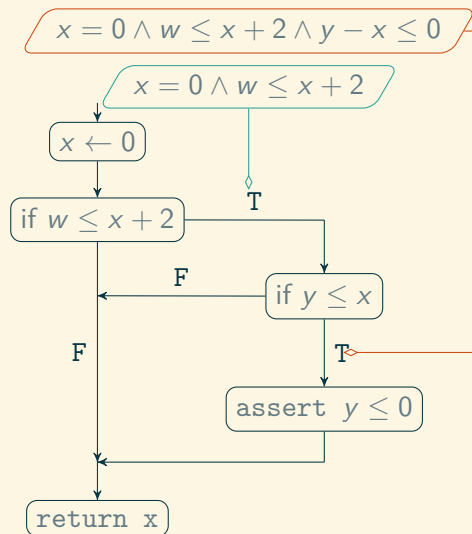
Static analysis computes invariants

Inequalities as invariants for a simple program

```

1 int example(int w, int y) {
2   int x = 0;
3   if (w <= x + 2) {
4     if (y <= x) {
5       assert y <= 0;
6     }
7   }
8   return x;
9 }

```



- Given some code, we produce a CFG
- Using DFA, we compute some invariants, for example
 - These two trapezoids represent the in-state and the out-state for the inner if
- Let's focus on this invariant, and describe briefly how Zones work

Zone Domain

$$x - Z_0 \leq 0$$

$$Z_0 - x \leq 0$$

$$w - x \leq 2$$

$$y - x \leq 0$$

$$y \leq 0$$

$$w \leq 2$$

Zonal state representation of data-flow analysis invariant

- Static program analysis computes invariants. We can represent these invariants with different ways. A popular choice for representing these invariants is the zonal domain as introduced by Miné.

Zone Domain

$$x - Z_0 \leq 0$$


$$Z_0 - x \leq 0$$

$$w - x \leq 2$$

$$y - x \leq 0$$

$$y \leq 0$$

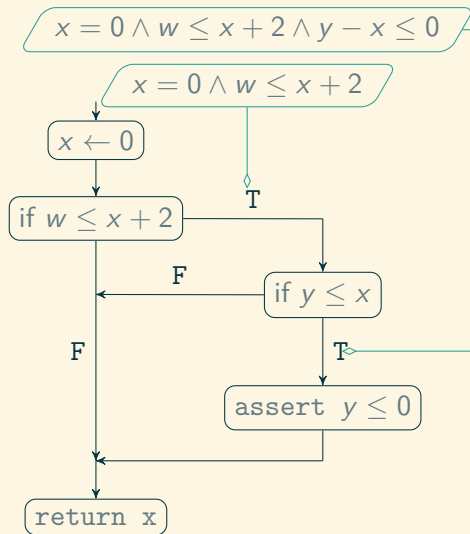
$$w \leq 2$$



- Static program analysis computes invariants. We can represent these invariants with different ways. A popular choice for representing these invariants is the zonal domain as introduced by Miné.

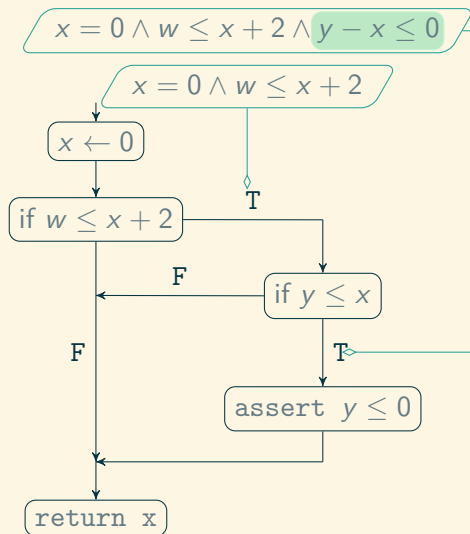
Zonal state representation of data-flow analysis invariant

Data-flow analysis incrementally updates variables



- Consider this simple data flow graph
- Using the intervals domain, for example, we know that the x variable is the only variable changed at each point along the true path.
- Using Zones, we can gain information about y .
- However, aside from small examples, it's unlikely that an update affects the entire state.

Data-flow analysis incrementally updates variables



- Consider this simple data flow graph
- Using the intervals domain, for example, we know that the x variable is the only variable changed at each point along the true path.
- Using Zones, we can gain information about y .
- However, aside from small examples, it's unlikely that an update affects the entire state.

Finding Affected Inequalities

Problem Definition


$$\begin{array}{rcl} x = 0 & & x = 0 \\ w - x \leq 2 & & w - x \leq 2 \\ y - x \leq 0 & \Rightarrow & y - x \leq 0 \\ \hline w \leq 2 & & \hline y \leq 0 \\ & & w \leq 2 \end{array}$$

What are the changed set of inequalities?

- Our problem is finding the smallest set of inequalities between two
- We approach this with several algorithms, which we now present.

Finding Affected Inequalities

Problem Definition




What are the changed set of inequalities?

- Our problem is finding the smallest set of inequalities between two
- We approach this with several algorithms, which we now present.

Finding Affected Inequalities


Problem Definition



What are the changed set of inequalities?

- Our problem is finding the smallest set of inequalities between two
- We approach this with several algorithms, which we now present.


Spurious Connected Variables¹



- Through inference, we may gain connections to variables which do not have an intentional logical connection
- Borrowing from Larsen et al. work, we remove these “spurious connections” as a pre-processing step to the remainder of our algorithms

¹Larsen et al., “Efficient Verification of Real-Time Systems: Compact Data Structure and State-Space Reduction”.


Connected Components



- Considering the state without the spurious connections
- The approach proposed by Visser et al. considers the undirected variable relation projection of the different constraints

Connected Components


Variable Relation Projection



- Considering the state without the spurious connections
- The approach proposed by Visser et al. considers the undirected variable relation projection of the different constraints

Connected Components


Variable Relation Projection



- Considering the state without the spurious connections
- The approach proposed by Visser et al. considers the undirected variable relation projection of the different constraints

Connected Components


Variable Relation Projection with impassable Z_0



- Considering the state without the spurious connections
- The approach proposed by Visser et al. considers the undirected variable relation projection of the different constraints

Connected Components


Variable Relation Projection with impassable Z_0



- Considering the state without the spurious connections
- The approach proposed by Visser et al. considers the undirected variable relation projection of the different constraints

Node Neighbors


Reconsider the out-going state without closed edges



- To demonstrate our next technique, we consider the same state without closure.
- The next technique extracts the forward-reachable and backward-reachable neighborhood of each changed variable.

Node Neighbors


Reconsider the out-going state without closed edges



- To demonstrate our next technique, we consider the same state without closure.
- The next technique extracts the forward-reachable and backward-reachable neighborhood of each changed variable.

Minimal Neighbors


Again, reconsider the out-going state without closed edges.



- Consider, instead, the edge of the updated edge.

Minimal Neighbors

Again, reconsider the out-going state without closed edges.



- Consider, instead, the edge of the updated edge.

Logically comparing different abstract domains

Research Questions

RQ1 Do the minimization algorithms reduce the size of a Zone state and improve runtime of domain comparisons?

RQ2 Do the minimization algorithms affect categorization of domain comparison results?

- We consider the following research questions for our application of empirically comparing abstract domains
- We compare Zones and Intervals, two comparable domains; and we compare Zones and Predicates, two incomparable domains

Experimental Setup

- Benchmarks: 127 Java methods
 - Ranging from 4 to 412 Jimple instructions
- Compared Zones to Intervals and Zones to Predicates
- Compared Total Runtime of Z3 to perform logical entailment of every combination, averaging over 5 executions

- Benchmark: 127 Java Methods selected from previous research
- Subject programs range from 4 to 412 Jimple instructions.

Experimental results show significant reduction in required number of inequalities for comparison

Average percentage changes in V and E between each technique

State Type	vs.	↓ Δ % V	↓ Δ % E
DFA Subject Programs			
CC	FS	70.37	29.47
NN	CC	0.02	0.01
MN	NN	0.10	0.05
EqBench Subject Programs			
CC	FS	43.0	2.1
NN	CC	0.0	0.0
MN	NN	0.13	0.13

- We see differences between the two benchmark suites due to the different qualities between the suites themselves
 - For example, the DFA benchmark suite is pulled from open-source projects with a high-preponderance of integer operations
 - Conversely, the EqBench consists of simply methods mainly used for equivalence testing within symbolic execution, many operations are not necessarily integer based

Experimental results show significantly reduced time to solver queries

State Type	~ Inter, sec.	~ Pred, sec.
DFA Subject Programs		
FS	4.03	265.91
CC	1.41	4.09
NN	1.41	4.04
MN	1.35	4.05
EQBench Subject Programs		
FS	0.79	5.56
CC	0.63	0.87
NN	0.58	0.9
MN	0.58	0.9

Experimental results show significant improvement in comparison granularity

State	\succ Intervals	= Intervals
DFA Subject Programs		
FS	2898	1002
CC	1194	2706
NN	1191	2709
MN	1164	2736
EQBench Subject Programs		
FS	374	255
CC	131	498
NN	131	498
MN	131	498

Experimental results show significant improvement in comparison granularity

State	\succ Predicates	$=$ Predicates	\prec Predicates	$\prec \succ$ Predicates
DFA Subject Programs				
FS	1464	237	167	2032
CC	1324	1930	473	173
NN	1322	1933	473	172
MN	1305	1960	473	162
EQBench Subject Programs				
FS	307	135	46	141
CC	217	322	72	18
NNy	217	322	72	18
MN	217	322	72	18

Conclusion

Experimental Results

- Minimization leads to reduced overall execution time when determining domain categorization.
- Minimization leads to improved granularity when evaluating domain precision.

Conclusion

Experimental Results

- Minimization leads to reduced overall execution time when determining domain categorization.
- Minimization leads to improved granularity when evaluating domain precision.

Algorithms and Approaches

- Spurious Connections → Reduce variable clustering
- Connected Components → Extract subsets using relational projection
- Node Neighbors → Extract subsets based on reachable neighborhoods
- Minimal Neighbors → Extract subsets leveraging semantic information

Future Work

- Extend to other Weakly-Relational Domains, e.g., Octagons
- Extend for comparison between relational domains

- Extend minimal identification to other weakly-relational domains, e.g., Octagons
- Extend minimal changes to a minimal union between two states for comparison of two weakly-relational domains.

Thank you


Questions?

The work reported here was supported by the U.S. National Science Foundation under award CCF-19-42044.

References I

- [1] K.G. Larsen et al. “Efficient Verification of Real-Time Systems: Compact Data Structure and State-Space Reduction”. In: *Proceedings Real-Time Systems Symposium*. IEEE Comput. Soc, 1997, pp. 14–24. ISBN: 081868268X. DOI: 10.1109/real.1997.641265.

Extended Examples of the Minimal Neighbors Algorithm



Extended Examples of the Minimal Neighbors Algorithm

